

MACHINE LEARNING ALGORITHMS

ALGORITHMS LIST :

1. Regression Algorithms

- Linear Regression
- Ridge Regression
- Lasso Regression

2. Dimensionality Reduction Algorithms

- Principal Component Analysis (PCA)
- Linear Discriminant Analysis (LDA)

3. Classification Algorithms

- Support Vector Machine (SVM)
- Decision Tree
- Random Forest

4. Clustering Algorithms

- K-Means Clustering
- Spectral Clustering
- Hierarchical Clustering

5. Deep Learning Algorithms

- Convolutional Neural Network (CNN)

1. REGRESSION ALGORITHMS

LINEAR REGRESSION

1. Initialize parameters:
 $w = 0$ // weight
 $b = 0$ // bias
 learning_rate = α
 max_iterations = N
2. Repeat until convergence or for max_iterations times:
 - a. Compute the predicted values:
 $y_{\text{pred}} = w * x + b$
 - b. Calculate the cost function (Mean Squared Error):
 $J = (1 / m) * \sum ((y_{\text{pred}}[i] - y[i])^2)$ for $i = 1$ to m
 - c. Compute the gradients:
 $dw = (2 / m) * \sum ((y_{\text{pred}}[i] - y[i]) * x[i])$ for $i = 1$ to m
 $db = (2 / m) * \sum (y_{\text{pred}}[i] - y[i])$ for $i = 1$ to m
 - d. Update the parameters using Gradient Descent:
 $w = w - \alpha * dw$
 $b = b - \alpha * db$
3. End Repeat
4. Output the final parameters:
 Print "Final weight (w): ", w
 Print "Final bias (b): ", b
5. Predict new values:
 $y_{\text{new}} = w * x_{\text{new}} + b$
 Return y_{new}

RIDGE REGRESSION

1. Initialize parameters:

```
w = 0 // weights  
b = 0 // bias  
learning_rate =  $\alpha$   
regularization_parameter =  $\lambda$   
max_iterations = N
```

2. Repeat until convergence or for max_iterations times:

a. Compute the predicted values:

```
y_pred = w * x + b
```

b. Calculate the cost function (Regularized Mean Squared Error):

```
J = (1 / m) *  $\sum (y\_pred[i] - y[i])^2$  +  $\lambda * \sum (w^2)$  for i = 1 to m
```

c. Compute the gradients:

```
dw = (2 / m) *  $\sum ((y\_pred[i] - y[i]) * x[i])$  + 2 *  $\lambda * w$   
db = (2 / m) *  $\sum (y\_pred[i] - y[i])$ 
```

d. Update the parameters:

```
w = w -  $\alpha * dw$   
b = b -  $\alpha * db$ 
```

3. End Repeat

4. Output the final parameters:

```
Print "Final weights (w): ", w  
Print "Final bias (b): ", b
```

5. Predict new values:

```
y_new = w * x_new + b  
Return y_new
```

LASSO REGRESSION

1. Initialize parameters:

```
w = 0 // weights
b = 0 // bias
learning_rate =  $\alpha$ 
regularization_parameter =  $\lambda$ 
max_iterations = N
```

2. Repeat until convergence or for max_iterations times:

a. Compute the predicted values:

```
y_pred = w * x + b
```

b. Calculate the cost function (Regularized Mean Squared Error):

```
J = (1 / m) *  $\sum (y\_pred[i] - y[i])^2$  +  $\lambda * \sum (|w|)$  for i = 1 to m
```

c. Compute the gradients:

```
dw = (2 / m) *  $\sum (y\_pred[i] - y[i]) * x[i]$ 
```

Adjust dw for L1 regularization:

```
if w > 0: dw +=  $\lambda$ 
```

```
if w < 0: dw -=  $\lambda$ 
```

```
db = (2 / m) *  $\sum (y\_pred[i] - y[i])$ 
```

d. Update the parameters:

```
w = w -  $\alpha * dw$ 
```

```
b = b -  $\alpha * db$ 
```

3. End Repeat

4. Output the final parameters:

```
Print "Final weights (w): ", w
```

```
Print "Final bias (b): ", b
```

5. Predict new values:

```
y_new = w * x_new + b
```

```
Return y_new
```

2.DIMENSIONALITY REDUCTION ALGORITHMS

PRINCIPAL COMPONENT ANALYSIS (PCA)

1. Input:

- Dataset X of size $(m \times n)$, where m is the number of samples and n is the number of features.

2. Preprocessing:

a. Standardize the dataset:

- For each feature j :
 - i. Compute mean μ_j and standard deviation σ_j .
 - ii. Standardize: $X_{\text{standardized}}[i][j] = (X[i][j] - \mu_j) / \sigma_j$

3. Compute the Covariance Matrix:

- $\text{Cov_matrix} = (1 / m) * (X_{\text{standardized}}^T * X_{\text{standardized}})$

4. Compute Eigenvalues and Eigenvectors:

- Find the eigenvalues and eigenvectors of Cov_matrix .

5. Sort Eigenvalues and Eigenvectors:

- Sort the eigenvalues in descending order.
- Arrange eigenvectors accordingly.

6. Select Principal Components:

- Choose the top k eigenvectors corresponding to the k largest eigenvalues.

7. Project Data:

- Project the standardized data onto the selected eigenvectors:
 $X_{\text{reduced}} = X_{\text{standardized}} * \text{Eigenvectors}_{\text{selected}}$

8. Output:

- Reduced dataset X_{reduced} of size $(m \times k)$.

LINEAR DISCRIMINANT ANALYSIS (LDA)

1. Input:

- Dataset X of size $(m \times n)$ and class labels y of size (m) .

2. Compute Class Statistics:

- For each class c :
 - a. Compute the mean vector μ_c for all features in class c .
 - b. Compute the overall mean vector μ for all samples.

3. Compute Within-Class Scatter Matrix (S_W):

- Initialize $S_W = 0$
- For each class c :
 - a. Compute the scatter matrix for class c :
$$S_{c} = \sum (x_i - \mu_c) * (x_i - \mu_c)^T \text{ for all } x_i \text{ in class } c$$
 - b. Add to S_W : $S_W += S_c$

4. Compute Between-Class Scatter Matrix (S_B):

- Initialize $S_B = 0$
- For each class c :
 - a. Compute $S_B += N_c * (\mu_c - \mu) * (\mu_c - \mu)^T$
where N_c is the number of samples in class c .

5. Solve Generalized Eigenvalue Problem:

- Solve: $S_B * w = \lambda * S_W * w$
- Find eigenvalues (λ) and eigenvectors (w).

6. Sort Eigenvalues and Eigenvectors:

- Sort eigenvalues in descending order.
- Arrange eigenvectors accordingly.

7. Select Discriminant Directions:

- Choose the top k eigenvectors corresponding to the k largest eigenvalues.

8. Project Data:

- Project the original data onto the selected eigenvectors:
$$X_{\text{reduced}} = X * \text{Eigenvectors}_{\text{selected}}$$

9. Output:

- Reduced dataset X_{reduced} of size $(m \times k)$.

3. CLASSIFICATION ALGORITHMS

SUPPORT VECTOR MACHINE (SVM)

1. Input:

- Dataset X of size $(m \times n)$ and labels y of size (m) , where $y \in \{+1, -1\}$.
- Regularization parameter C (for Soft Margin SVM).
- Learning rate α .
- Maximum iterations N .

2. Initialize:

- Initialize weights $w = 0$ (size n).
- Initialize bias $b = 0$.

3. Repeat until convergence or for N iterations:

a. For each training example (x_i, y_i) :

i. Compute the decision function:

$$f(x_i) = (w \cdot x_i) + b$$

ii. Check the classification condition:

$$\text{If } y_i * f(x_i) \geq 1:$$

- Update weights with no penalty:

$$w = w - \alpha * \lambda * w$$

Else:

- Update weights and bias with penalty:

$$w = w - \alpha * (\lambda * w - y_i * x_i)$$

$$b = b + \alpha * y_i$$

iii. $\lambda = (1 / C)$ for soft margin SVM, $\lambda = 0$ for hard margin SVM.

4. Output:

- Final weight vector w and bias b .

5. Prediction:

- For a new input x , predict:

$$y_{\text{pred}} = \text{sign}((w \cdot x) + b)$$

DECISION TREE

TRAINING PHASE

1. Input:

- Dataset D with features X and target variable y.
- Stopping criteria (e.g., max depth, minimum samples per node).

2. Define:

- Impurity measure:
 - For classification: Entropy (ID3) or Gini Index (CART).
 - For regression: Mean Squared Error (MSE).
- Splitting criteria:
 - Choose the feature and threshold that minimize impurity or maximize information gain.

3. Function: BuildTree(D, depth)

a. Check for stopping criteria:

- If all examples in D belong to the same class, return a leaf node with the class label.
- If depth exceeds the max depth or D has fewer than the minimum samples, return a leaf node with the majority class (classification) or mean value (regression).

b. Find the best split:

i. For each feature x_j in X:

- For each possible threshold t:
 - Split D into left (D_{left}) and right (D_{right}) subsets:
 $D_{left} = \{\text{examples in } D \text{ where } x_j \leq t\}$
 $D_{right} = \{\text{examples in } D \text{ where } x_j > t\}$
 - Compute impurity for the split (weighted by subset sizes):
$$\text{Impurity} = (|D_{left}| / |D|) * \text{Impurity}(D_{left}) + (|D_{right}| / |D|) * \text{Impurity}(D_{right})$$
- ii. Select the feature x_j and threshold t with the lowest impurity.

c. Create a decision node:

- Feature: x_j
- Threshold: t

d. Recursively build left and right subtrees:

$\text{left_subtree} = \text{BuildTree}(D_{left}, \text{depth} + 1)$
 $\text{right_subtree} = \text{BuildTree}(D_{right}, \text{depth} + 1)$

e. Return the decision node with:

- Feature: x_j
- Threshold: t
- Left Subtree: left_subtree
- Right Subtree: right_subtree

4. Call BuildTree(D, 0) to build the tree.

PREDICTION PHASE:

1. Input:

- Decision tree T.
- New input example x.

2. Traverse the tree:

- a. Start at the root node.
- b. At each decision node:
 - Check the feature and threshold:
If $x[\text{feature}] \leq \text{threshold}$:
Move to the left subtree.
 - Else:
Move to the right subtree.

3. When a leaf node is reached:

- Return the class label (classification) or mean value (regression) stored in the leaf.

4. Output:

- Predicted value or class for x.

RANDOM FOREST

TRAINING PHASE

1. Input:
 - Dataset D with features X and target variable y.
 - Number of trees T.
 - Number of features to select for each tree f (default: $\sqrt{\text{total features}}$).
 - Maximum tree depth, minimum samples per leaf, or other stopping criteria.
2. Initialize:
 - Forest = [] // List to store all decision trees.
3. For t = 1 to T:
 - a. Bootstrap Sampling:
 - Create a bootstrap sample D_t by randomly sampling with replacement from D.
 - b. Feature Subset Selection:
 - Randomly select f features from the total features.
 - c. Train a Decision Tree:
 - BuildTree(D_t, selected_features) using the decision tree algorithm.
 - Use only the selected f features for splitting at each node.
 - d. Add the trained tree to the Forest:
 - Forest.append(tree)
4. Output:
 - Trained random forest (Forest containing T trees).

PREDICTION PHASE:

1. Input:
 - Trained Random Forest (Forest).
 - New input example x.
2. Initialize:
 - Predictions = [] // To store predictions from all trees.
3. For each tree in the Forest:
 - a. Traverse the tree to predict the value or class for x:
 - prediction = tree.predict(x)
 - b. Append prediction to Predictions.
4. Aggregate Predictions:
 - For classification:
 - Final_prediction = Majority vote from Predictions.
 - For regression:
 - Final_prediction = Average of Predictions.
5. Output:
 - Final_prediction (class or value).

4. CLUSTERING ALGORITHMS

K-MEANS CLUSTERING

1. Input:

- Dataset D with n data points $\{x_1, x_2, \dots, x_n\}$.
- Number of clusters k .
- Maximum number of iterations max_iter .
- Tolerance value tol for convergence.

2. Initialize:

- Randomly select k points from D as initial centroids $\{c_1, c_2, \dots, c_k\}$.
- Set iteration counter $\text{iter} = 0$.

3. Repeat until convergence or max_iter is reached:

- Assign each data point to the nearest centroid:
 - For each data point x_i in D :
Assign x_i to cluster j such that:
 $j = \text{argmin}(\|x_i - c_j\|^2) \text{ for all } j \text{ in } \{1, 2, \dots, k\}$.
- Update centroids:
 - For each cluster j in $\{1, 2, \dots, k\}$:
Update centroid $c_j = \text{Mean of all points assigned to cluster } j$.
- Check for convergence:
 - If the change in centroids ($\|c_{\text{old}} - c_{\text{new}}\|$) is less than tol for all centroids, break.
- Increment iteration counter: $\text{iter} = \text{iter} + 1$.

4. Output:

- Final cluster centroids $\{c_1, c_2, \dots, c_k\}$.
- Cluster assignments for all data points.

SPECTRAL CLUSTERING

1. Input:

- Dataset D with n data points $\{x_1, x_2, \dots, x_n\}$.
- Number of clusters k .

2. Construct the Similarity Matrix:

- Compute a similarity matrix S of size $(n \times n)$:
 - For each pair of points (x_i, x_j) :
 $S[i][j] = \text{similarity}(x_i, x_j)$
(e.g., Gaussian similarity: $S[i][j] = \exp(-||x_i - x_j||^2 / (2\sigma^2))$).
- Set diagonal elements $S[i][i] = 0$ (self-similarity is not considered).

3. Construct the Graph Laplacian:

- Compute the degree matrix D :
 - $D[i][i] = \sum S[i][j]$ for all j .
- Compute the unnormalized Laplacian L :
 - $L = D - S$.
- (Optional) Normalize the Laplacian:
 - $L_{\text{sym}} = D^{-1/2} * L * D^{-1/2}$ (symmetric normalization).
 - $L_{\text{rw}} = D^{-1} * L$ (random walk normalization).

4. Compute Eigenvalues and Eigenvectors:

- Compute the first k eigenvectors of L (or $L_{\text{sym}}/L_{\text{rw}}$) corresponding to the k smallest eigenvalues.

5. Form the Feature Matrix:

- Construct a matrix U of size $(n \times k)$, where the columns are the k eigenvectors.

6. Normalize the Rows of U (Optional):

- For each row i in U , normalize it to have unit length:
 $U[i] = U[i] / ||U[i]||$.

7. Apply K-Means Clustering:

- Treat each row of U as a point in k -dimensional space.
- Use the K-Means algorithm to cluster the rows into k clusters.

8. Output:

- Cluster assignments for all data points in D .

HIERARCHICAL CLUSTERING

1. Input:

- Dataset D with n data points $\{x_1, x_2, \dots, x_n\}$.
- Linkage criterion (e.g., single linkage, complete linkage, average linkage, or ward linkage).
- Desired number of clusters k or stopping condition (e.g., a distance threshold).

2. Initialize:

- Each data point is its own cluster. So, initially, there are n clusters.
- Create a distance matrix to store pairwise distances between all points.
 - For each pair of points (x_i, x_j) , compute the distance $d(x_i, x_j)$.
 - Typically, Euclidean distance is used:
$$d(x_i, x_j) = ||x_i - x_j||.$$

3. Repeat until there are k clusters or no further merging possible:

- Find the two clusters (C_i, C_j) with the smallest distance (based on the chosen linkage criterion).
 - For **single linkage**: Use the minimum pairwise distance between clusters.
 - For **complete linkage**: Use the maximum pairwise distance between clusters.
 - For **average linkage**: Use the average pairwise distance between clusters.
 - For **Ward's linkage**: Minimize the total variance of the merged clusters.
- Merge the two closest clusters C_i and C_j into a single cluster C_{ij} .
 - Remove C_i and C_j from the list of clusters.
 - Add the new merged cluster C_{ij} to the list of clusters.
- Update the distance matrix:
 - Recalculate the distances from the newly formed cluster C_{ij} to all other clusters.
 - If using **single linkage**, update the distance to the minimum distance between C_{ij} and any other cluster.
 - If using **complete linkage**, update the distance to the maximum distance between C_{ij} and any other cluster.
 - If using **average linkage**, update the distance to the average distance between C_{ij} and any other cluster.

4. Output:

- Dendrogram: A tree structure representing the merging process of clusters.
- Cluster assignments (if stopping at a desired number of clusters k).

5. DEEP LEARNING ALGORITHMS

CONVOLUTIONAL NEURAL NETWORK (CNN)

1. Initialize parameters:
 - Define the number of layers, filters, filter sizes, stride, padding, and activation functions.
 - Initialize weights and biases for all filters and fully connected layers.
2. Input:
 - Provide the input image or feature map (e.g., size $H \times W \times C$, where H =height, W =width, C =channels).
3. Forward Propagation:
 - a. Convolutional Layer:
 - For each filter in the layer:
 - i. Slide the filter over the input (using stride and padding as defined).
 - ii. Perform element-wise multiplication between the filter and the input region.
 - iii. Sum the results to produce a single value.
 - iv. Add the bias term.
 - Generate an output feature map.
 - b. Apply Activation Function:
 - Apply the chosen activation function (e.g., ReLU) to each value in the feature map.
 - c. Pooling Layer (if applicable):
 - Divide the feature map into pooling regions.
 - Perform the pooling operation (e.g., max-pooling or average-pooling) within each region.
 - Reduce the spatial dimensions of the feature map.
 - d. Repeat steps 3a–3c for each convolutional and pooling layer.
 - e. Flattening:
 - Convert the final feature maps into a 1D vector.
 - f. Fully Connected Layers:
 - Multiply the flattened vector by the weights of the fully connected layer and add biases.
 - Apply the activation function.
4. Output Layer:
 - Use the final layer to output predictions:
 - For classification tasks, apply softmax for probabilities.
 - For regression tasks, output the raw values.
5. Backpropagation:
 - a. Compute Loss:
 - Calculate the loss using a suitable loss function (e.g., cross-entropy for classification, MSE for regression).

b. Compute Gradients:

- Use backpropagation to calculate the gradients of the loss with respect to weights and biases for all layers.

c. Update Parameters:

- Update weights and biases using Gradient Descent or an optimization algorithm like Adam:

$\text{weight} = \text{weight} - \text{learning_rate} * \text{gradient}$

$\text{bias} = \text{bias} - \text{learning_rate} * \text{gradient}$

6. Repeat:

- Iterate over multiple epochs and batches of the training data.

7. Prediction:

- Use the trained CNN to make predictions on new inputs.